

Effective Java, Chapter 4

Classes and Interfaces



- Material From Joshua Bloch
 - Effective Java: Programming Language Guide
- Cover Parts of Items 15 through 24
 - "Classes and Interfaces" Chapter
- Moral:
 - Inheritance requires careful programming



Item 15: Minimize Accessibility of Classes and Members

- Information hiding /encapsulation
 - Most important factor to distinguish well or poorly designed module
 - Hide the implementation detail
 - Modules are communicated through API



More on Item 15

- Reasons for Information Hiding
 - Decouples modules
 - Allows isolated development and maintenance
 - Developed/tested/optimized/used/understood /modified in isolation
- Java has an <u>access control mechanism</u> to accomplish this goal



Accessibility Levels

- Standard list of accessibility levels
 - private
 - package-private (aka package friendly)
 - protected
 - public
- Huge difference between 2nd and 3rd
 - package-private: part of implementation
 - protected: part of public API



Make Each Class or Member as Inaccessible as Possible

- For top-level classes/interfaces
 - Only package-private and public are applicable
 - Always favor package-private if possible
- If a package-private top-level class/interface is used by only one class
 - Consider making it a *private* nested class
- Choose the <u>less inaccessible level</u> if you are not sure



More on access level

- Accessibility of methods which overrides a superclass method
 - Not permitted to have a lower access level.
 - Note: all members of an interface are implicitly public, thus its overridden method in subclass must be public.
- Don't raise the level of accessibility just to facilitate testing



Module System in Java 9

- A module is a grouping of packages
- A module may explicitly export some of its packages via export declarations in its module declaration
 - which is by convention contained in a source file named module-info.java
- Public and protected members of public classes in unexported packages give rise to the two implicit access levels



Item 16 Use Accessors, Not Public Fields

Avoid code such as:

```
Public Class Person { public String name; public int age; }
```

- No possibility of encapsulation
- No chance to check validity
- public mutable fields are **not** thread safe

Use get/set methods instead:

```
public double getX() { return x; }
public void setX(double x) { this.x = x}
```

- Advice holds for immutable fields as well
 - Limits possible ways for class to evolve

Example: Questionable Use of Immutable Public Fields

```
public final class Time {
   private static final int HOURS PER DAY = 24;
  private static final int MINUTES PER HOUR = 60;
  public final int hour; // Not possible to change rep for class
  public final int minute;
   // But we can check invariants, since fields are final
   public Time ( int hour, int minute ) {
      if (hour < 0 | hour >= HOURS PER DAY)
         throw new IllegalArgumentException ("Hour: " + hour);
      if (minute < 0 | minute >= MINUTES PER HOUR)
         throw new IllegalArgumentException("Minute: " + minute);
      this.hour = hour;
      this.minute = minute;
   // Remainder omitted
```

Code example: violation of this rule in Java platform lib

```
package java.awt;
// ...
public class Dimension extends Dimension2D
    implements java.io.Serializable {
    /**
    * The width dimension; negative values can be used.
    public int width:
    /**
    * The height dimension; negative values can be used.
   public int height;
   // ...
   public Dimension getSize() {
   return new Dimension (width, height);
```

Not possible to make the fields immutable later.



Exception: Exposing Constants

- Public constants must contain either <u>primitive values</u> or references to <u>immutable objects</u>
- Wrong Potential Security Hole:

```
public static final Type[] VALUES = {...};
```

- Problem:
 - values is final; entries in values are not!



Exposing Constants - Solutions

Correct:

```
private static final Type[] PRIVATE_VALUES = {...};
public static final List VALUES =
    Collections.unmodifiableList(Arrays.asList
    (PRIVATE_VALUES));
```

Also Correct:

```
private static final Type[] PRIVATE_VALUES = {...};
public static final Type[] values() {
   return (Type[]) PRIVATE_VALUES.clone();
}
```



Item 17: Minimize Mutability

- An object is considered <u>immutable</u> if its state cannot change after it is constructed.
 - e.g String, Integer, BigInteger, BigDecimal
- Reasons to make a class immutable:
 - Thread safe inherently
 - Can be shared freely
 - 3. No need to make defensive copy
 - Its internals can also be shared
 - 5. Great building blocks for other object
 - Great to be map keys and set elements

Code example: a mutable class

```
oublic class SynchronizedRGB {
   private int red;
   private int green;
   private int blue;
   private String name;
   public SynchronizedRGB (int red.
                           int green,
                           int blue,
                           String name) {
       this.red = red:
       this.green = green;
       this.blue = blue;
       this.name = name;
   public void set(int red,
                    int green,
                    int blue,
                    String name) {
       synchronized (this) {
           this.red = red:
           this.green = green;
           this.blue = blue;
           this.name = name;
```

The class, <u>SynchronizedRGB</u>, defines objects that represent colors. Each object represents the color as three integers that stand for primary color values and a string that gives the name of the color.

```
public synchronized int getRGB() {
    return ((red << 16) | (green << 8) | blue);
}

public synchronized String getName() {
    return name;
}

public synchronized void invert() {
    red = 255 - red;
    green = 255 - green;
    blue = 255 - blue;
    name = "Inverse of " + name;
}</pre>
```

Use of a mutable class should be very careful in concurrent application

```
synchronized (color) {
   int myColorInt = color.getRGB();
   String myColorName = color.getName();
}
```

The SynchronizedRGB object may be seen in an inconsistent state:

If another thread invokes color.set after Statement 1 but before Statement 2, the value of myColorInt won't match the value of myColorName.

To avoid this outcome, the two statements must be bound together.

This kind of inconsistency is only possible for <u>mutable objects</u> — it will not be an issue for the <u>immutable</u> ones.



Strategy for Defining Immutable Objects

- Don't provide any <u>mutators</u> ("setter" methods)
- 2) Make all fields *final* and *private*
- 3) Don't allow subclasses to *override* methods.
- If the instance fields include <u>references to</u> <u>mutable objects</u>, don't allow those objects to be changed:
 - Don't provide methods that modify the mutable objects.
 - Don't share references to the mutable objects.

Code example: immutable version

```
final public class ImmutableRGB {
   // Values must be between 0 and 255.
   private final int red;
   private final int green;
   private final int blue;
   private final String name;
   public ImmutableRGB (int red,
                        int green,
                        int blue,
                        String name) {
        this.red = red:
        this.green = green;
        this.blue = blue;
        this.name = name;
```

ImmutableRGB.java

Immutable objects are inherently thread-safe; they require no synchronization



Disadvantage: Performance

- Typical approach:
 - Provide immutable class
 - Provide <u>mutable companion class</u> for situations where performance is an issue
- Clients choose on performance needs
- Example in Java Library:
 - String (Immutable)
 - StringBuilder (Companion Mutable Class)
 - StringBuffer (Deprecated Companion Mutable Class)
- Static factories can cache frequently used items



More discussions: should all fields of immutable class be marked as final?

Why String class is immutable even though it has a non -final field called "hash"



I was reading through Item 15 of Effective Java by Joshua Bloch. Inside Item 15 which speaks about 'minimizing mutability' he mentions five rules to make objects immutable. One of them is is to make all fields final. Here is the rule:



Make all fields final: This clearly expresses your intent in a manner that is enforced by the system. Also, it is necessary to ensure correct behavior if a reference to a newly created instance is passed from one thread to another without synchronization, as spelled out in the memory model [JLS, 17.5; Goetz06 16].

I know that String class is an example of a immutable class. Going through the source code I see that it actually has a hash instance which is not final.

```
//Cache the hash code for the string
private int hash; // Default to 0
```

How does String become immutable then?



More discussions: should all fields of immutable class be marked as final?

- "No methods may be modify the object and that all its fields must be final" is <u>a bit</u> <u>stronger than necessary</u>.
- In truth, no method may produce an "<u>external visible</u>" change in the object's state.
 - The hash field in String is used to cache the hash code for the string, see detail explanation at stackoverflow





- Interface inheritance does NOT have these problems
- Interface inheritance is better for lots of reasons...

Inheritance breaks encapsulation!

- In other words, a subclass depends on the <u>implementation details</u> of its superclass for its proper function (see example in next few slides)
- Difficult to evolve superclass without breaking subclasses

Example: Broken Subtype

```
// This is very common, but broken
public class InstrumentedHashSet<E> extends HashSet<E>
  private int addCount = 0; // add() calls
  public InstrumentedHashSet() {}
 public InstrumentedHashSet(Collection<? extends E> c)
     { super(c); }
  public boolean add(E o) {
     addCount++; return super.add(o);
  public boolean addAll(Collection<? extends E> c) {
     addCount += c.size(); return super.addAll(c);
  // accessor method to get the count
  public int addCount() { return addCount; }
```



Broken Example, continued

So, what's the problem?



```
InstrumentedHashSet<String> s =
    new InstrumentedHashSet<String>();
s.addAll(Arrays.asList("Snap", "Crackle", "Pop"));
```

- What does addCount() return?
 - **3**?
 - **6**?
- Internally, HashSet addAll() is implemented on top of add(), which is overridden. Note that this is an implementation detail, so we can't get it right in InstrumentedHashSet.



- Overriding methods can be tricky
 - May break the superclass invariant
 - Overridden method does not maintain invariant
 - May break subclass invariant
 - New methods may be added to superclass
 - What if new method matches subclass method?
- Also, recall problems with equals() and hashCode()



Composition

 Fortunately, composition solves all of these problems – even though it makes the programmer work harder

```
// Inheritance
public Class Foo extends File {...}

// Composition
public class Foo {
   private File f = ...;
   // Note: forwarded(转发) methods
...
}
```

Revisiting the Example

```
public class InstrumentedSet<E> implements Set<E> {
   private final Set <E> s;
   private int addCount = 0;
   public InstrumentedSet (Set<E> s) { this.s = s}
   public boolean add(E o) {
      addCount++; return s.add(o); }
   public boolean addAll (Collection<? extends E> c) {
      addCount += c.size();
      return s.addAll(c);
   // forwarded methods from Set interface
```



Note that an InstrumentedSet IS-A Set

```
public class InstrumentedSet<E> implements Set<E> {
   private final Set <E> s;
   private int addCount = 0;
   public InstrumentedSet (Set<E> s) { this.s = s}
   public boolean add(E o) {
      addCount++; return s.add(o); }
   public boolean addAll (Collection<? extends E> c) {
      addCount += c.size();
      return s.addAll(c);
   // forwarded methods from Set interface
```



- Consider temporarily instrumenting a Set
- Note that Set is an interface

```
Set<Date> s = new InstrumentedSet<Date>(new TreeSet<Date>(cmp))
Set<E> s2 = new InstrumentedSet<E>(new HashSet<E>(capacity));
static void f(Set<Dog> s) {
    InstrumentedSet<Dog> myS = new InstrumentedSet<Dog> (s);
    // use myS instead of s
    // all changes are reflected in s!
}
```

Violations in Java platform libs

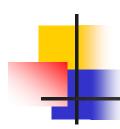
There are a number of obvious violations of this principle in the Java platform libraries.

```
public
class Stack<E> extends Vector<E> {
public
class Properties extends Hashtable<Object,Object> {
```

- In case of <u>Properties</u>, the designer intended that only strings be allowed as keys and values
 - Properties.setProperty only accepts only strings
 - But Hashtable.put can accept any object
 - Invariant of Properties is violated and it was too late to corret ...



- Or else prohibit it.
 - Mark the class as <u>final</u>
 - Make the constructor <u>private</u>
- Document effects of overriding any method
 - Document "self use" of overridable methods
 - This is "implementation detail", but unavoidable, since subclass "sees" implementation.
- Inheritance violates encapsulation!



Code example

java.util.AbstractCollection

```
public boolean remove (Object o)
```

Removes a single instance of the specified element from this collection, if it is present (optional operation). More formally, removes an element e such that (o==null ? e==null : o.equals(e)), if this collection contains one or more such elements. Returns true if this collection contained the specified element (or equivalently, if this collection changed as a result of the call).

This implementation iterates over the collection looking for the specified element. If it finds the element, it removes the element from the collection using the iterator's remove method.

Note that this implementation throws an UnsupportedOperationException if the iterator returned by this collection's iterator method does not implement the remove method and this collection contacts the specified object.

```
public boolean remove(Object o) {
  Iterator<E> e = iterator();
  if (o==null) {
    while (e.hasNext()) {
      if (e.next()==null) {
          e.remove();
          return true;
      }
    }
} else {
    while (e.hasNext()) {
    if (o.equals(e.next())) {
        e.remove();
        return true;
    }
    }
} return false;
}
```

The description reveals the implementation detail...



Inheritance is Forever

- A commitment to allow inheritance is <u>part of</u> <u>public API</u>
 - If you provide a poor interface, you (and all of the subclasses) are stuck with it.
 - You cannot change the interface in subsequent releases.
- Write subclass to test a class designed for inheritance
 - To find whether a method should be made <u>protected</u>
 - Usually three subclasses are sufficient

Constructors Must Not Invoke Overridable Methods

```
// Problem - constructor invokes overridden m()
public class Super {
  public Super() { m();}
  public void m() {...};
}
public class Sub extends Super {
  private final Date date;
  public Sub() {date = new Date();}
  public void m() { // access date variable}
}
```



What Is the Problem?

Consider the code

```
Sub s = new Sub();
```

- The first thing that happens in Sub() constructor is a call to constructor in Super()
- The call to m() in Super() is overridden
- But date variable is not yet initialized!
- 4) Further, initialization in Super m() never happens!
- Yuk!



Inheritance and Cloneable, Serializable

Since clone() and readObject() behave a lot like constructors, these methods cannot invoke overridable methods either.

Problem – access to uninitialized state



Bottom Line

- Be sure you want inheritance, and design for it
- Don't inherit from ordinary concrete class consider:
 - Make class <u>final</u>,
 - Make constructors <u>private</u> (or package-private)



Item 20 Interfaces & Abstract Classes

- Existing classes can be easily retrofitted(改 进,翻新) to <u>implement a new interface</u>
 - Same is not true for abstract classes due to single inheritance model in Java
- Interfaces are ideal for "mixins"
 - Example: Comparable interface
- Interfaces aren't required to form a <u>hierarchy</u>
 - Some things aren't hierarchical



More on interface

- But interface can not contain method implementation
 - Sometimes too tedious to use
 - You have to implement a dozen methods before you can create an object

```
add(E): boolean
   add(int, E): void
   addAll(int, Collection <? extends E>): boolean
♠ addAll(Collection<? extends E>): boolean
a clear(): void
a contains(Object) : boolean
⊕ a containsAll(Collection <?>): boolean
♠ a equals(Object) : boolean
   get(int) : E
a hashCode(): int
   indexOf(Object): int
isEmpty() : boolean
iterator() : Iterator < E >
   lastIndexOf(Object): int
   listIterator(): ListIterator<E>
   listIterator(int): ListIterator < E >
   remove(int) : E
remove(Object) : boolean
removeAll(Collection<?>): boolean
set(int, E) : E
size(): int
   subList(int, int) : List < E >
```



Advantage of Abstract Classes

- It is far <u>easier to evolve</u> an abstract class than an interface
 - It is relatively easy to <u>add a new method</u> to an abstract class
 - But almost impossible to add a method to a public interface
 - Should be very carefully to design public interfaces



- Combine the <u>virtues of interfaces & abstract</u> classes
 - The <u>interface</u> still defines the type
 - But the <u>skeletal implementation</u> takes all of the work out of implementing it
- By convention, skeletal implementation are named <u>AbstractInterface</u>
 - E.g. AbstractCollection, AbstractSet, AbstractList, and AbstractMap



```
// Concrete implementation built on abstract skeleton
static List<Integer> intArrayAsList(final int[] a) {
   if (a == null) throw new NPE(...);
   // Note anonymous class
   return new AbstractList() {
     public Object get(int i) {
        return new Integer(a[i]);
      public int size() { return a.length; }
      public Object set(int i, Object o) {
         int oldVal = a[i];
         a[i] = ((Integer) o).intValue();
         return new Integer (oldVal);
   };
```

This Implementation Does a Lot!

- The List interface includes many methods.
- Only 3 of them are explicitly provided here.
- This is an "anonymous class" example
 - Certain methods are overridden



- Favor interface over abstract class.
 - Unless <u>ease of evolution</u> is more important than flexibility and power.
- Do understand the <u>limitations of abstract class</u>.
- Skeletal implementation can be considered to go with a nontrivial interface.
- Design <u>public interface</u> carefully and do through test before publishing.

Item 23: Prefer Class Hierarchies to Tagged Classes

```
//Tagged Class - vastly inferior to a class hierarchy
class Figure {
                                                  Add new shape is not easy
  enum Shape { RECTANGLE, CIRCLE };
  final Shape shape; // Tag field
  double length; double width; // for RECTANGLE
  double radius;
                                // for CIRCLE
  Figure (double length, double width) {...} // RECTANGLE
  Figure (double radius) {...} // CIRCLE
  double area() {
       switch (shape) { // Gag! Roll-your-own dispatching!
          case RECTANGLE: return length*width;
          case CIRCLE: return Math.PI*(radius * radius);
          default: throw new AssertionError();
```

More on Item 20

```
//Tagged Class - vastly inferior to a class hierarchy
class Figure {
                                               Add new shape is not easy
  enum Shape { RECTANGLE, CIRCLE };
  final Shape shape; // Tag field
                                                 Fields not applicable
  double length; double width; // for RECTANGLE
                                                 to all Figures and
  double radius;
                              // for CIRCLE
                                                 can not make final
  Figure (double length, double width) {...} // RECTANGLE
  Figure (double radius) {...} // CIRCLE
  double area() {
      case RECTANGLE: return length*width;
          case CIRCLE: return Math.PI*(radius * radius);
          default: throw new AssertionError();
                                               Multiple implementation
```



- Define an abstract class
 - <u>Abstract method</u> each method in the tagged class whose behavior depends on the tag value
 - E.g. *area()* in *Figure* class
 - <u>Final method</u> each method whose behavior does not depend on the tag value
 - Common fields shared by all the specific types
- Define a concrete class for each specific type
 - Appropriate implementation for each abstract method.
 - Data fields particular to this type

A much better solution

```
//Class hierarchy replacement for a tagged class
abstract double area();
class Circle extends Figure {
  final double radius;
  Circle(double rad) { radius = rad; }
  double area() {
       return Math.PI * (radius * radius);
class Rectangle extends Figure {
  final double length; final double width;
  Rectangle (double len; double wid)
     { length = len; width = wid; }
  double area() {
     return length * width;
```



Item 24: Favor Static Member Classes Over Nonstatic

- Nested classes (嵌套类)
 - Defined within another class
 - Used to serve its enclosing class (外围类)
 - Inner class methods can <u>access the private data</u> from the scope in which they are defined.
 - Hidden from other classes in the same package.
 - Anonymous inner classes are handy to define callbacks without writing a lot of code



Four flavors of nested classes

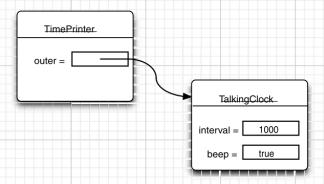
- Static member class
- Nonstatic member class (inner)
- Anonymous class (inner)
- Local class (inner)

See code examples to understand the differences.



Static vs NonStatic

- Static <u>requires no connection</u> to enclosing instance.
- Nonstatic always <u>associated with enclosing</u> instance
 - Possible performance issue
 - Possible desire to create by itself



Hence recommendation to favor static



Local classes

- Declared anywhere a local variable may be declared
- Same scoping rules
- Have names like member classes
- May be static or nonstatic (depending on whether context is static or nonstatic)



More facts about nested class

- Inner classes are a phenomenon of the <u>compiler</u>, not the <u>virtual machine</u>.
- Local/anonymous class can access the outer

local final variable.

The beep in the class
 TimePrinter is a copy.
 (Refer to CJ1 6.4.5 for detail)

```
public void start(int interval, final boolean beep)
{
   class TimePrinter implements ActionListener
   {
      public void actionPerformed(ActionEvent event)
      {
            Date now = new Date();
            System.out.println("At the tone, the time is " + now);
            if (beep) Toolkit.getDefaultToolkit().beep();
      }
   }
}

ActionListener listener = new TimePrinter();
   Timer t = new Timer(interval, listener);
   t.start();
}
```